

High-Level Design (HLD) Document

The HLD provides an architectural overview, defining the main components and their interactions in the machine learning system.

1. System Goal

To predict the **Cryptocurrency Liquidity Index** using historical market data, enabling proactive detection of potential liquidity crises and informing risk management strategies.

2. Architectural Diagram

The system follows a standard batch-processing Machine Learning Pipeline architecture.

3. Key Components and Technologies

Component	Description	Technologies
Data Source	Raw historical cryptocurrency market data (price, volume, market cap).	CSV files (coin_gecko_*.csv)
Data Ingestion	Module for loading, merging, and initial cleaning of the raw data files.	Pandas
Data Processor	Handles data quality checks, missing value imputation, and feature scaling.	Pandas, NumPy, Scikit-learn (MinMaxScaler)
Feature Engineer	Creates the target Liquidity Index and other derivative features (e.g., Volume/Price Ratio).	Pandas, NumPy
ML Model	The core prediction engine selected for non-linear regression and robustness.	Scikit-learn (RandomForestRegressor)
Model Evaluation	Measures model performance against key business and statistical metrics.	Scikit-learn (RMSE, MAE, R^2)
Persistence Layer	Stores the trained model and the feature scaler for later use/deployment.	Joblib (.pkl files)

4. Data Flow

1. **Ingestion:** Raw .csv files are read and merged into a single DataFrame.
2. **Preprocessing:** Data is cleaned (median imputation) and prepared for feature creation.
3. **Feature Engineering:** The **Liquidity_Index** target is calculated.
4. **Scaling & Split:** All numerical features are scaled, and the dataset is split into training (80%) and testing (20%) sets.
5. **Training:** The **RandomForestRegressor** is trained on the scaled training data.
6. **Evaluation:** The model's predictions on the test set are evaluated using defined metrics.
7. **Output:** The final model and scaler are saved to disk using joblib.

Low-Level Design (LLD) Document

The LLD details the implementation of each core module, including function specifications, inputs, outputs, and algorithms.

1. Project Directory Structure

```
├── cryptocurrency_liquidity_prediction.py # Main ML Pipeline script
├── coin_gecko_2022-03-17.csv           # Data File 1
├── coin_gecko_2022-03-16.csv           # Data File 2
├── liquidity_predictor.pkl              # Output: Trained Model
└── data_scaler.pkl                     # Output: Trained Scaler
```

2. Module Specifications (Function Detail)

A. *load_and_merge_data(file_paths)*

- **Purpose:** Reads raw market data from specified CSV files and concatenates them.
- **Algorithm:** Iterative read, pd.concat, and sorting by coin and date.
- **Inputs:** file_paths (List of strings, e.g., ['file1.csv', 'file2.csv']).
- **Outputs:** df_combined (Pandas DataFrame).

B. *preprocess_data(df)*

- **Purpose:** Cleans the DataFrame by handling non-numeric columns and imputing missing values.
- **Algorithm:**
 1. Drop categorical columns (coin, symbol, date).
 2. Use the **Median** to fill all remaining NaN values across all numerical columns.
- **Inputs:** df (Raw Pandas DataFrame).
- **Outputs:** df_cleaned (Pandas DataFrame with only numerical, imputed features).

C. *feature_engineering(df)*

- **Purpose:** Creates the predictive target variable and a highly correlated feature.
- **Algorithm:**
 - Target Creation (Liquidity_Index):
$$\text{Liquidity Index} = \frac{\text{24h_volume}}{\text{mkt_cap}} \cdot |\text{24h change}| + \epsilon$$
 - Feature Creation (Vol_Price_Ratio):
$$\text{Vol_Price_Ratio} = \frac{\text{24h_volume}}{\text{price} + \epsilon}$$
- **Inputs:** df (Cleaned DataFrame).
- **Outputs:** df_featured (DataFrame including the Liquidity_Index target).

D. *prepare_for_training(df)*

- **Purpose:** Scales features and splits the data into training and testing sets.
- **Algorithm:**

1. Initialize **MinMaxScaler**.
 2. Fit and transform features (X).
 3. Use **train_test_split** with `test_size=0.2` and `shuffle=False` (to maintain time order).
- **Inputs:** df (Featured DataFrame).
 - **Outputs:** X_train, X_test, y_train, y_test, scaler (fitted MinMaxScaler).

E. *train_and_evaluate_model(X_train, X_test, y_train, y_test)*

- **Purpose:** Trains the chosen model and assesses its performance.
- **Algorithm:**
 1. **Model:** Instantiate **RandomForestRegressor** (`n_estimators=100`, `max_depth=10`).
 2. **Training:** Call `model.fit(X_train, y_train)`.
 3. **Prediction:** Call `model.predict(X_test)`.
 4. **Evaluation Metrics:** Calculate **RMSE**, **MAE**, and **R^2 score**.
- **Inputs:** Train/Test splits (X and y).
- **Outputs:** model (Trained **RandomForestRegressor** object).

3. Model Selection Justification

Approach	Model	Justification
Algorithm	RandomForestRegressor	Non-linear model capable of capturing complex interactions between price, volume, and market cap. Robust against high data skewness and less sensitive to feature collinearity compared to linear models.
Scaling	MinMaxScaler	Essential for normalizing the vast differences in magnitude between features (e.g., price vs. mkt_cap) to ensure fair contribution during model training.